

Assignment 2 : MA20268

Disclaimer: This solution to assignments is made (typed & thought upon) completely by hand and with full understanding of the content by me

– Shubham Krishan, 24MA10063

Constraints: Using only `numpy` for calculations and `matplotlib` for visualizations

Question 1:

1. Generate two vectors $u \in \mathbb{R}^{200}$ and $v \in \mathbb{R}^{200}$ sampled independently from a uniform random distribution on the range $[-5, 5]$.

Define the target variable $y \in \mathbb{R}^{200}$ by

$$y_i = \begin{cases} 1, & \text{if } u_i^2 + v_i^2 > 4, \\ 0, & \text{otherwise,} \end{cases} \quad i = 1, \dots, 200. \quad (1)$$

Here, y_i , u_i , and v_i denote the i -th entries of the vectors y , u , and v , respectively.

Now, define the following basis functions:

$$\begin{aligned} g_0(u_i, v_i) &= 1, & g_1(u_i, v_i) &= u_i, & g_2(u_i, v_i) &= v_i, \\ g_3(u_i, v_i) &= u_i^2, & g_4(u_i, v_i) &= v_i^2, & g_5(u_i, v_i) &= u_i v_i. \end{aligned} \quad (2)$$

Use the basis functions to define the matrix $C \in \mathbb{R}^{200 \times 6}$ where

$$C_i^j = g_j(u_i, v_i), \quad (3)$$

with i as the row index and j as the column index.

Use the least-squares method to predict the target variable y by learning the coefficient vector $w \in \mathbb{R}^6$ in the linear model

$$y \approx Cw. \quad (4)$$

Figure 1: Q1

Initializing $u, v \in \mathbb{R}^{200}$

```
import numpy as np

u = np.random.uniform(-5, 5, 200)
v = np.random.uniform(-5, 5, 200)
```

Now, the target variable $y \in \mathbb{R}^{200}$ as defined:

```
y = np.where(u**2 + v**2 > 4, 1, 0)
```

Given basis functions: $g_0(u_i, v_i) = 1$, $g_1(u_i, v_i) = u_i$, $g_2(u_i, v_i) = v_i$, $g_3(u_i, v_i) = u_i^2$, $g_4(u_i, v_i) = v_i^2$, $g_5(u_i, v_i) = v_i u_i$

```
def g1(ui, vi):
    return 1

def g2(ui, vi):
    return ui

def g3(ui, vi):
    return vi

def g4(ui, vi):
    return ui**2

def g5(ui, vi):
    return vi**2

def g6(ui, vi):
    return ui*vi

basis_func = {
    0: g1,
    1: g2,
    2: g3,
    3: g4,
    4: g5,
    5: g6
}
```

Defining matrix $C \in \mathbb{R}^{200 \times 6}$

```
c = np.zeros((200, 6))

for i in range(200):
    for j in range(6):
        c[i][j] = basis_func.get(j, lambda u,v: None)(u[i], v[i])
```

Training for $\tilde{y}$:

the model for this problem is:

$$\tilde{y}(u, v) = \sum_{i=0}^5 w_i g_i(u, v) = w_0 + w_1 u + w_2 v + w_3 u^2 + w_4 v^2 + w_5 uv$$

That is, $\mathbf{y} = C\mathbf{w}$

to use least squares method, define: $E(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^{200} (y_i - \tilde{y}_i)^2$

Now, our main goal is to minimize $E(\mathbf{w})$.

to get $\tilde{w}$ we solve the least squares problem $\min_w \|y - Cw\|^2$

```
w = np.linalg.lstsq(c, y, rcond=None)[0]
```

```
print("Weights: ")
for i, val in enumerate(w):
    print(i+1, val)
```

```
Weights:
1 0.7023864289928273
2 0.002929697069766947
3 -0.00045019473195849556
4 0.01308140541652499
5 0.01145844422715964
6 0.0016065574587148636
```

Final views on the problem

The target set and basis functions used here trained the model $\tilde{y}$ to predict whether a point (u, v) lies outside the circle of radius '2' or not (defined by the target condition: $u^2 + v^2 > 4$). So, now we should be able to predict this using the approximated weights $\tilde{w}$.

```
def pred(u, v, w):
    features = np.array([1,u,v,u**2,v**2,u*v])
    y_hat = features @ w # dot product of given surface with w
    return y_hat

TEST_SIZE = 100
MAX_PRINT = 10
u_test = np.random.uniform(-5, 5, TEST_SIZE)
v_test = np.random.uniform(-5, 5, TEST_SIZE)
print("TESTING on random uniform test with size: ", TEST_SIZE )

correct = 0
threshold = np.mean(y)

for i in range(TEST_SIZE):
    y_hat = pred(u_test[i], v_test[i], w)
    prediction = 1 if y_hat > threshold else 0
    t = 1 if u_test[i]**2 + v_test[i]**2 > 4 else 0

    if prediction == t:
        correct += 1
    if i < MAX_PRINT:
        print(f"({u_test[i]:.2f}, {v_test[i]:.2f}) => model: {prediction} truth: {t}")
    elif i == MAX_PRINT:
        print("...")

print(f"points outside the circle in training data: {threshold}%") # shows skewness in the
↪ random training dataset
print(f"accuracy: {(correct/TEST_SIZE)*100}%")
```

```
TESTING on random uniform test with size: 100
(-4.14, 2.39) => model: 1 truth: 1
(-0.92, -3.16) => model: 0 truth: 1
(4.39, -0.69) => model: 1 truth: 1
(-3.74, 2.01) => model: 1 truth: 1
(-0.42, 1.29) => model: 0 truth: 0
```

```

(4.85, -0.77) => model: 1 truth: 1
(2.02, 3.01) => model: 0 truth: 1
(-1.00, 0.62) => model: 0 truth: 0
(3.55, 1.46) => model: 1 truth: 1
(-4.10, 0.67) => model: 1 truth: 1
...
points outside the circle in training data: 0.905%
accuracy: 62.0%

```

Note: Since the dataset is not balanced (i.e., the proportion of class 1 labels is significantly higher than that of class 0), a threshold of 0.5 is not appropriate. As the model is obtained via least-squares regression and produces continuous outputs, a decision threshold is required for classification. In this case, we use the empirical mean of the training labels as the threshold to account for the class imbalance.

Question 2:

- Download the Fashion-MNIST dataset (both train and test sets, 28×28 images). Crop each image by 4 pixels from both rows and columns (resulting in $20 \times 20 = 400$ dimensions per vectorized image).

Let x_i be a sample in the train set and $c(x_i)$ be its true class label (0–9). Run 9 separate simulations comparing class 0 (T-shirt/top) against each of the other classes, where:

- $c(x_i) = 0 \Rightarrow \text{label } +1$

1

- $c(x_i) = k \ (k \neq 0) \Rightarrow \text{label } -1$

The 9 experiments are:

$$\begin{aligned}
 &0 \text{ vs } 1, \quad 0 \text{ vs } 2, \quad 0 \text{ vs } 3, \\
 &0 \text{ vs } 4, \quad 0 \text{ vs } 5, \quad 0 \text{ vs } 6, \\
 &0 \text{ vs } 7, \quad 0 \text{ vs } 8, \quad 0 \text{ vs } 9
 \end{aligned} \tag{5}$$

For each experiment, use least squares to train a linear regression model $f(\cdot)$ on the binary train subset.

On the test set, predict:

$$\hat{c}(x_j) = \begin{cases} +1 & \text{if } f(x_j) > 0 \\ -1 & \text{otherwise} \end{cases} \tag{6}$$

Using ground truth test labels, compute the misclassification error for samples truly belonging to class 0 or the opposing class k .

Setting up MNIST dataset

using keras datasets for clean mnist import to avoid messy mnist_reader util files (alternative way was to use zalando's mnist_reader util)

```
from tensorflow.keras.datasets import fashion_mnist
(x_train, y_train), (x_test, y_test) = fashion_mnist.load_data()
```

To reduce dimensionality and remove irrelevant border pixels, we crop image from 28×28 to 20×20 :

```
x_train = x_train[:, 4:24, 4:24]
x_test = x_test[:, 4:24, 4:24]
```

Each image is flattened into a feature vector:

```
x_train = x_train.reshape(len(x_train), -1)
x_test = x_test.reshape(len(x_test), -1)

# normalize pixel values to [0,1] to improve numerical stability of least squares
x_train = x_train / 255.0
x_test = x_test / 255.0
```

Setting up the Model

First, we define a helper function to *extract only the relevant classes (0 and k) and relabel them*:

```
# Helper function to filter out x, y sets into two label classes 0 and k
def get_binary(x, y, k):
    mask = (y==0) | (y==k)
    x_bin = x[mask]
    y_bin = y[mask]

    y_bin = np.where(y_bin == 0, 1, -1);

    return x_bin, y_bin
```

Model: Least Square Classifier

We model the prediction as: $\tilde{y} = \mathbf{w}^T \mathbf{x} + b$, where b is the bias term.

To incorporate the bias term, we augment the feature matrix with a column of ones. The optimal weight $\tilde{w}$ are obtained by solving the equation:

$$\min_{\mathbf{w}} |\mathbf{X}\mathbf{w} - \mathbf{y}|^2$$

Using the Moore-Penrose pseudoinverse (built-in in numpy):

```
def train_ls(X, y):
    X = np.c_[np.ones(X.shape[0]), X] # add bias
    w = np.linalg.pinv(X) @ y # best fit line using pseudoinverse
    return w
```

Prediction Rule:

Predictions are made using the sign of the model output:

```
def predict(X, w):
    X = np.c_[np.ones(X.shape[0]), X]
    return np.where(X @ w > 0, 1, -1)
```

Evaluation Metric:

We use the misclassification error defined as: $\text{Error} = \frac{1}{N} \sum_{i=1}^N \mathbb{1}(y_i \neq \hat{y}_i)$

```
def misclassification_error(y_true, y_pred):  
    return np.mean(y_true != y_pred)
```

Running Experiments

Training and evaluating the model for each $k \in 1, \dots, 9$:

```
results = {}  
  
for k in range(1, 10):  
    X_train_k, y_train_k = get_binary(x_train, y_train, k)  
    w = train_ls(X_train_k, y_train_k)  
    X_test_k, y_test_k = get_binary(x_test, y_test, k)  
    y_pred = predict(X_test_k, w)  
    error = misclassification_error(y_test_k, y_pred)  
  
    results[k] = error  
    print(f"0 vs {k} => error: {error:.4f}")
```

```
0 vs 1 => error: 0.0205  
0 vs 2 => error: 0.0515  
0 vs 3 => error: 0.0730  
0 vs 4 => error: 0.0300  
0 vs 5 => error: 0.0050  
0 vs 6 => error: 0.1750  
0 vs 7 => error: 0.0005  
0 vs 8 => error: 0.0400  
0 vs 9 => error: 0.0025
```

Interpretation and Inference

This experiment highlights how well a linear model trained via least squares can separate different classes in a high-dimensional image space.

Key observations:

- Since the model is linear, performance depends on how linearly separable the classes are.
- Some class pairs (e.g., visually distinct clothing types) are expected to yield lower error.
- Others with similar shapes or textures may result in higher misclassification.

Despite its simplicity, this approach provides a useful baseline before moving to more expressive models like logistic regression or neural networks.

Question 3:

3. Generate two vectors $u \in \mathbb{R}^{100}$ and $v \in \mathbb{R}^{100}$ sampled independently from a uniform random distribution on the range $[-5, 5]$.

Define the target variable $y \in \mathbb{R}^{100}$ by

$$y_i = \begin{cases} +1 & \text{if } u_i v_i > 0.5, \\ -1 & \text{otherwise,} \end{cases} \quad (7)$$

where y_i, u_i, v_i denote the i -th entries of y, u, v , respectively.

Now define the following basis functions:

$$\begin{aligned} g_0(u_i, v_i) &= 1, & g_1(u_i, v_i) &= u_i, & g_2(u_i, v_i) &= v_i, \\ g_3(u_i, v_i) &= u_i^2, & g_4(u_i, v_i) &= v_i^2, & g_5(u_i, v_i) &= u_i v_i. \end{aligned} \quad (8)$$

Use these to form the matrix $C \in \mathbb{R}^{100 \times 6}$ with entries $C_i^j = g_j(u_i, v_i)$, where i is the row index and j is the column index.

Use least squares to learn a coefficient vector $w \in \mathbb{R}^6$ in the linear model $y \approx Cw$. Let $\hat{y} = Cw$ be the predictions, and define the Mean Squared Error (MSE) as

$$\text{MSE} = \sum_{i=1}^{100} (y_i - \hat{y}_i)^2. \quad (9)$$

Then compute MSE.

Figure 2: Q3

Process

Our goal is to find MSE for the target variable $y \in \mathbb{R}^{100}$ given by:

$$y_i = \begin{cases} +1 & \text{if } u_i v_i > 0.5 \\ -1 & \text{otherwise} \end{cases}$$

i.e. a prediction model to show if a point $(u, v) \in \{\text{set of points outside hyperbola } uv = 0.5 \text{ for } u \in \mathbf{u}, v \in \mathbf{v}\}$

The dataset for training is *random points sampled independently from the uniform distribution* of points in a square of side length 5.

Defining the Training Dataset

```
Q3_SAMPLE_SIZE = 100
u = np.random.uniform(-5, 5, Q3_SAMPLE_SIZE)
v = np.random.uniform(-5, 5, Q3_SAMPLE_SIZE)
```

and the target variable y is:

```
y = np.where(u*v > 0.5, 1, -1)
```

Building the $\mathbb{C}$ matrix with given basis functions;

```
c = np.zeros((Q3_SAMPLE_SIZE, 6))

# using the basis function dictionary defined in Q1
for i in range(Q3_SAMPLE_SIZE):
    for j in range(6):
        c[i][j] = basis_func.get(j, lambda u,v: None)(u[i], v[i])
```

Solving for the optimal weights

The optimal weights $\mathbf{w}$ such that $\mathbb{C}\mathbf{w} \approx \mathbf{y}$ is calculated as:

```
w = np.linalg.lstsq(c, y, rcond=None)[0]

print("weights: ")
for i, val in enumerate(w):
    print(f"{i} => {val}")
```

```
weights:
0 => -0.17947304158961683
1 => -0.023836947181572907
2 => 0.016270735557239027
3 => 0.0178503760140826
4 => -0.004682018450994885
5 => 0.07458843435837348
```

Calculating MSE

Here, we have $\hat{y} = \mathbb{C}\mathbf{w}$ and MSE is given by:

$$\text{MSE} = \sum_{i=1}^{100} (y_i - \hat{y}_i)^2$$

```
y_hat = c @ w
mse = np.sum((y-y_hat)**2)

print(f"mse = {mse:.5f}, mean: {np.mean((y-y_hat)**2)}")
```

```
mse = 42.54981, mean: 0.42549809815599815
```


Question 4:

4. Generate a vector $r \in \mathbb{R}^{100}$ sampled from a uniform random distribution on $[-8, 8]$. Define the target variable $t \in \mathbb{R}^{100}$ by

$$t_i = e^{0.1r_i} \cos(r_i), \quad (10)$$

where t_i, r_i denote the i -th entries of t, r .

Define the first set of basis functions:

$$\begin{aligned} h_0(r_i) &= 1, & h_1(r_i) &= r_i, & h_2(r_i) &= r_i^2, \\ h_3(r_i) &= r_i^3, & h_4(r_i) &= r_i^4, & h_5(r_i) &= r_i^5. \end{aligned} \quad (11)$$

Form the matrix $B \in \mathbb{R}^{100 \times 6}$ with $B_i^j = h_j(r_i)$. Use least squares to fit coefficients $\beta \in \mathbb{R}^6$ minimizing $\|B\beta - t\|_2^2$, and compute

$$\text{MSE}_1 = \frac{1}{100} \sum_{i=1}^{100} (t_i - \hat{t}_i)^2. \quad (12)$$

Now add three higher-order basis functions:

$$h_6(r_i) = r_i^6, \quad h_7(r_i) = r_i^7, \quad h_8(r_i) = r_i^8, \quad (13)$$

forming $B' \in \mathbb{R}^{100 \times 9}$ and coefficients $\beta' \in \mathbb{R}^9$. Compute MSE_2 similarly.

What happens to the MSE when adding the three new basis functions? (Hint: Plot predictions vs. ground truth to understand overfitting behavior.)

Figure 3: Q4

Inference and Observations

- The model approximates a nonlinear boundary ($uv = 0.5$) using linear least squares with basis functions.
- Performance depends on whether the basis includes interaction terms (e.g., $\mathbf{uv}$).
- A lower MSE indicates better approximation of the nonlinear relationship.
- This is effectively a regression-based classification, where predictions are thresholded.

New Things Learned:

Python code:

- `np.linalg.lstsq(A, b, rcond=None)` => solves the matrix linear equation directly: $A\mathbf{x} = b$
- `np.where(<condition>, <val_T>, <val_F>)` => allowed me to construct np vectors directly using conditions

Inference of problems:

Question 1:

- understood the differences in applications between regression and classification models
- got more familiar with terminologies like class labels

- understood the use of model and basis function in the QUERY $\Rightarrow$ PREDICTION flow.

Question 2:

- Kuchh to seekha hu lekin likhna bhul gya $\hat{_}$ ^'